

Mandelbrot Set Image Generation

Maths, Data, and Machines

Gareth Lloyd

ACCU York

2nd July 2025

Outline

- 1 Introduction
- 2 Generating the Image
- 3 Optimization Journey
- 4 Conclusion

What is the Mandelbrot Set?

- A famous mathematical set of complex numbers
- Boundary forms a fractal pattern
- High-quality visualization made by Benoit Mandelbrot in 1980
- First defined and drawn by Robert W. Brooks and Peter Matelski in 1978
- Generated by iterating a simple mathematical formula

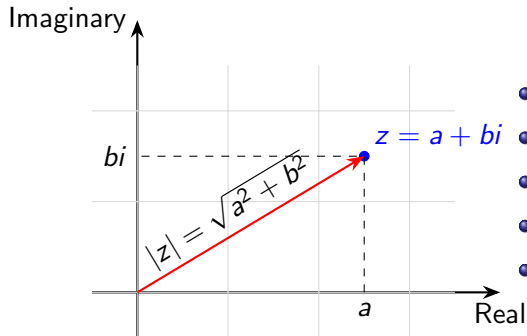
The Mathematical Foundation

The Mandelbrot Formula

$$z_{n+1} = z_n^2 + c$$

- Starting with $z_0 = 0$
- c is a point in the complex plane
- If the sequence remains bounded, the point is in the set
- A sequence diverges to infinity if $|z_n| > 2$, hence iteration can be stopped

Complex Numbers



- Complex number $z = a + bi$
- a : real part (x-axis)
- b : imaginary part (y-axis)
- $|z| = \sqrt{a^2 + b^2}$: magnitude
- $i^2 = -1$

Algorithm Overview

- ① Choose a region in the complex plane
- ② For each pixel:
 - ▶ Map to complex number c
 - ▶ Iterate the formula
 - ▶ Color based on escape time (number of iterations)

Mandelbrot Set Visualization

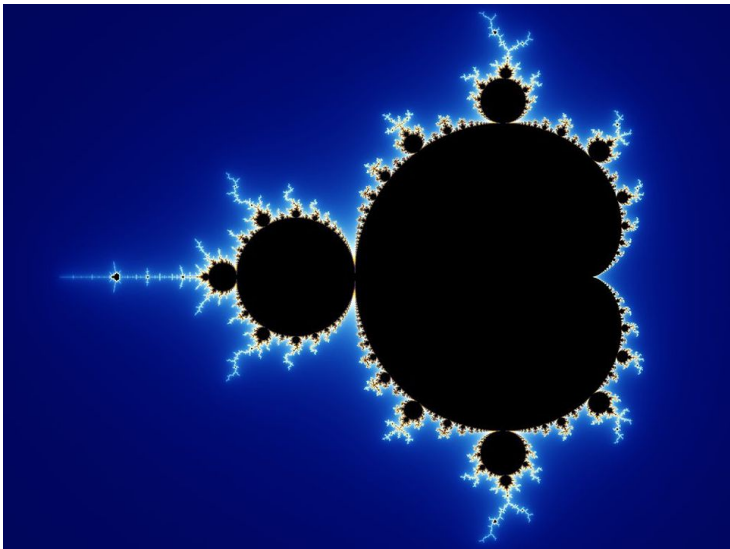


Image by Wolfgang Beyer (CC BY-SA 3.0)

Naïve

Using `std::complex`

- Clean, readable code using standard library
- `std::complex` handles complex arithmetic
- Easy to understand and maintain

Implementation

```
auto mandelbrot(std::complex<double> c) -> std::size_t {  
    auto iter = std::size_t{};  
    auto z = std::complex<double>{};  
  
    while (std::abs(z) <= 2.0 and iter < MAX_ITER) {  
        z = z * z + c;  
        ++iter;  
    }  
    return iter;  
}
```


Engineering Mindset

Evidence-Driven Development

- Avoid premature optimization based on assumptions
- Let evidence guide your efforts
- Similar to TDD: Know when something is wrong because the tests fail

Performance Optimization Cycle

- 1 Profile and measure current performance
- 2 Identify actual bottlenecks (not assumed ones)
- 3 Make targeted improvements
- 4 Verify with measurements
- 5 Repeat

Benchmarking Setup

Google Benchmark Code

```
static void BM_Mandelbrot_V1(benchmark::State &state) {  
    for (auto _ : state) {  
        auto result = mandelbrot<MAX_ITER>(std::complex{0.0, 0.0});  
        benchmark::DoNotOptimize(result);  
    }  
    state.counters["calc"] = benchmark::Counter(1, benchmark::Counter::kIsIterationInvariantRate);  
}  
BENCHMARK(BM_Mandelbrot_V1);
```

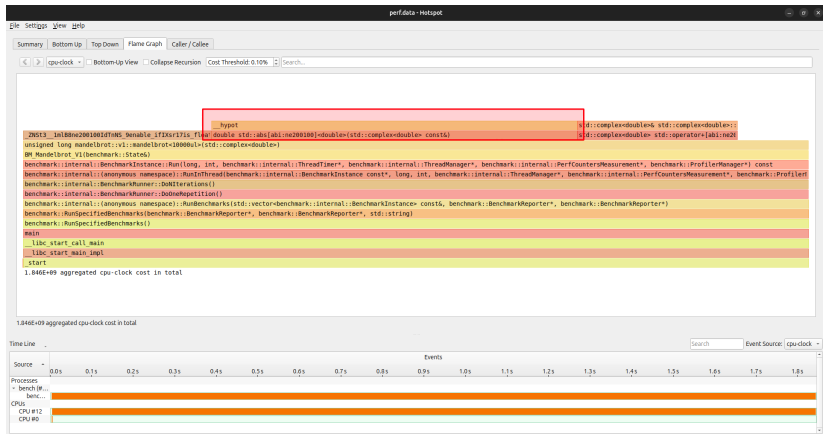
Profiling Tools

- Google Benchmark for microbenchmarking
- Linux perf for performance analysis
- Hotspot as a convenient GUI for perf

Why 0.0 + 0.0i?

- This is the worst case, it is in the set and can not escape
- Are there other cases we should profile?

Performance Analysis



- Profiling reveals significant time spent in `std::abs()`

The Problem with `std::abs`

Performance Issues

- `std::abs` for `std::complex` calls `__hypot`
- `__hypot` is a non-inlinable C library call
- `__hypot` internally performs an expensive square root operation

Mathematical Insight

$$\begin{aligned}\text{hypot}(z) \leq 2 &\Leftrightarrow \sqrt{a^2 + b^2} \leq 2 \\ &\Leftrightarrow a^2 + b^2 \leq 4 \\ &\Leftrightarrow \text{norm}(z) \leq 4\end{aligned}$$

Optimized

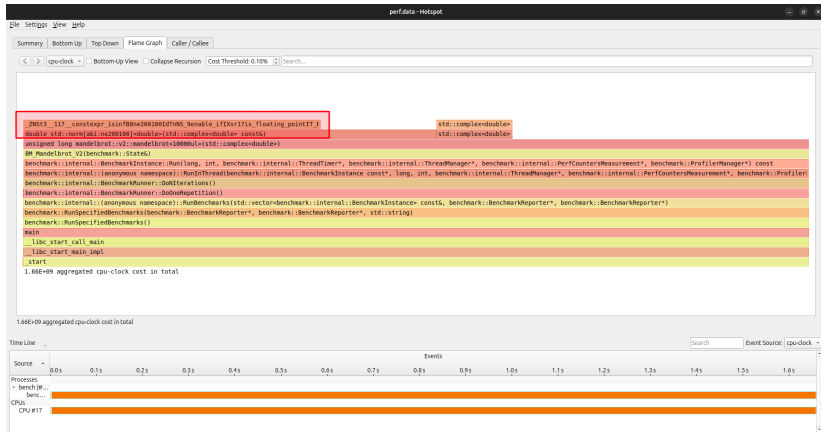
Using std::norm

- Replaces `std::abs(z) <= 2` with `std::norm(z) <= 4`
- Avoids expensive square root operation, maintains mathematical equivalence

Implementation

```
auto mandelbrot(std::complex<double> c) -> std::size_t {  
    auto iter = std::size_t{};  
    auto z = std::complex<double>{};  
  
    while (std::norm(z) <= 4.0 and iter < MAX_ITER) {  
        z = z * z + c;  
        ++iter;  
    }  
    return iter;  
}
```

Performance Analysis



- Small amount of work to handle infinities
- Our values stay small we don't need to handle infinities

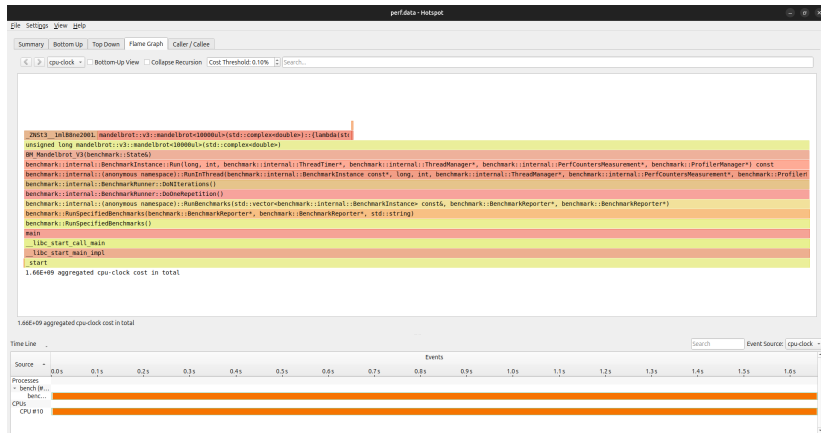
Avoiding `std::norm`

- Directly compute $x^2 + y^2$ instead of using `std::norm`
- No need to handle special cases (infinities, NaNs)
- More explicit control over the computation

Implementation

```
auto mandelbrot(std::complex<double> c) -> std::size_t {  
    auto iter = std::size_t{};  
    auto z = std::complex<double>{};  
  
    auto not_escaped = [](std::complex<double> z) {  
        auto x = z.real();  
        auto y = z.imag();  
        return x * x + y * y <= 4.0;  
    };  
  
    while (not_escaped(z) and iter < MAX_ITER) {  
        z = z * z + c;  
        ++iter;  
    }  
    return iter;  
}
```

Performance Analysis



- Both implementations have nearly identical performance
- `std::norm` templated code is fully inlined by the compiler

Removing Abstractions

Why Remove `std::complex`?

- Potential overhead from abstraction
- Want explicit control over operations
- Better optimization opportunities

Mathematical Derivation

Starting with $z = x + yi$ and $c = a + bi$:

$$\begin{aligned} z^2 + c &= (x + yi)^2 + (a + bi) \\ &= (x^2 + 2xyi + (yi)^2) + (a + bi) \\ &= (x^2 + 2xyi + y^2 i^2) + (a + bi) \\ &= (x^2 + 2xyi - y^2) + (a + bi) \quad \text{since } i^2 = -1 \\ &= x^2 + 2xyi - y^2 + a + bi \\ &= x^2 - y^2 + a + 2xyi + bi \\ &= (x^2 - y^2 + a) + i(2xy + b) \end{aligned}$$

No Abstraction

Implementation

```
auto mandelbrot(std::complex<double> c) -> std::size_t {  
    auto const a = c.real();  
    auto const b = c.imag();  
    auto iter = std::size_t{};  
    auto x = 0.0;  
    auto y = 0.0;  
  
    while (x * x + y * y <= 4.0 and iter < MAX_ITER) {  
        auto x_next = x * x - y * y + a;  
        auto y_next = 2 * x * y + b;  
        std::tie(x, y) = std::tie(x_next, y_next);  
        ++iter;  
    }  
    return iter;  
}
```

Details

- Uses `std::tie` for clean value updates
- Avoids temporary complex number objects
- Easier for compiler to see operation dependencies
- Better quality codegen, is faster

Computation Reuse

Key Insight / Experiment

- x^2 and y^2 are used in both the escape check and value update
- We can compute them once per iteration and reuse them
- Reduces redundant multiplications

Optimized Implementation

```
auto mandelbrot(std::complex<double> c) -> std::size_t {  
    auto const a = c.real();  
    auto const b = c.imag();  
    auto iter = std::size_t{};  
    auto x = 0.0;  
    auto y = 0.0;  
    auto x2 = 0.0; // x squared  
    auto y2 = 0.0; // y squared  
  
    while (x2 + y2 <= 4.0 and iter < MAX_ITER) {  
        auto x_next = x2 - y2 + a;  
        auto y_next = 2 * x * y + b;  
        std::tie(x, y) = std::tie(x_next, y_next);  
        y2 = y * y; // store to reuse in the loop check  
        x2 = x * x; // store to reuse in the loop check  
        ++iter;  
    }  
    return iter;  
}
```

Beyond Hotspot: Going Deeper with perf

Hotspot's Limits

- Hotspot was good for seeing the big picture, functions and lines of code
- Now we need instruction-level analysis
- Time to use Linux perf directly for detailed insights

Using perf assembly

```
# Record with call graph and debug info
perf record -g --call-graph=dwarf ./bench
--benchmark_filter=BM_Mandelbrot_V4/0
--benchmark_min_time=1s

# View annotated assembly
perf annotate
```

Assembly Comparison

	1e0:	nop
		mov %rcx,0x80(%rsp)
		dec %r14
	↓	je 24f
	1ed:	xorpd %xmm1,%xmm1
		xor %eax,%eax
		xorpd %xmm2,%xmm2
		nop
2.57	200:	movapd %xmm1,%xmm3
3.27		movapd %xmm2,%xmm4
5.11		mulsd %xmm2,%xmm4
5.60		mulsd %xmm1,%xmm1
5.06		subsd %xmm4,%xmm1
5.68		addsd %xmm5,%xmm1
6.21		addsd %xmm3,%xmm3
6.27		mulsd %xmm3,%xmm2
5.37		addsd %xmm6,%xmm2
5.25		lea 0x1(%rax),%rcx
5.48		movapd %xmm2,%xmm3
5.69		mulsd %xmm2,%xmm3
5.63		movapd %xmm1,%xmm4
5.60		mulsd %xmm1,%xmm4
5.44		addsd %xmm3,%xmm4
10.31		ucomisd %xmm4,%xmm0
3.18	↑	jb 1e0
3.01		cmp \$0x270f,%rax
2.43		mov %rcx,%rax
2.83	↑	jb 200
		jmp 1e0
	24f:	mov %rbx,%rdi

No Abstraction

	cs	nopw 0x0(%rax,%rax,1)
	1e0:	mov %rcx,0x80(%rsp)
		dec %r14
	↓	je 247
	1ed:	xorpd %xmm3,%xmm3
		xorpd %xmm2,%xmm2
		xor %eax,%eax
		xorpd %xmm4,%xmm4
		xorpd %xmm1,%xmm1
		nop
5.98	200:	subsd %xmm3,%xmm2
6.27		addsd %xmm4,%xmm4
7.03		mulsd %xmm4,%xmm1
7.71		movapd %xmm2,%xmm4
7.36		addsd %xmm6,%xmm4
8.27		addsd %xmm7,%xmm1
6.01		movapd %xmm1,%xmm3
5.85		mulsd %xmm1,%xmm3
4.75		movapd %xmm4,%xmm2
4.67		mulsd %xmm4,%xmm2
4.10		lea 0x1(%rax),%rcx
4.27		movapd %xmm2,%xmm5
3.98		addsd %xmm3,%xmm5
8.02		ucomisd %xmm5,%xmm0
3.69	↑	jb 1e0
3.49		cmp \$0x270f,%rax
4.09		mov %rcx,%rax
4.42	↑	jb 200
		jmp 1e0
0.03		mov %rbx,%rdi
	247:	

Computation Reuse

Assembly Comparison

Key Observations

- Manual reuse can still be a beneficial optimization
 - ▶ No abstraction: 4x moves, 5x multiplications, 6x add/subs
 - ▶ Computation reuse: 4x moves, 3x multiplications, 6x add/subs
- 2 fewer multiplications per iteration with reuse
- Manual optimization beneficial despite compiler optimizations

Modern x86_64 Microarchitecture Levels

Our Current Code: Using x86_64-v1 Instructions

- Using `mulsd` (Multiply Scalar Double)
- Part of SSE2 (Streaming SIMD Extensions 2)
- Operates on just 2 registers at a time

x86_64 Microarchitecture Levels

x86_64-v2 POPCNT, CMPXCHG16B

- Common since 2011 (Nehalem/Bulldozer)

x86_64-v3 AVX, AVX2, BMI1/2, FMA, LZCNT

- Common since 2015 (Haswell/Excavator)

Modern Compiler Targeting

CMake Configuration

```
target_compile_options(  
    bench PRIVATE -march=x86-64-v3 -mtune=native  
)
```

Why Target x86_64-v3?

- Realistic for today's deployments
- Enables vectorized operations (AVX/AVX2)
- Modern math instructions (FMA)
- Bit manipulation (BMI1/2)
- Better bit operations (LZCNT)

Assembly Comparison x86_64-v3

	data10	data10	CS	noplw	0x0(%rax)
1e0:	mov	%rcx,0x80(%rsp)			
	dec	%r14			
	je	23a			
1ed:	vxorpd	%xmm1,%xmm1,%xmm1			
	vxorpd	%xmm2,%xmm2,%xmm2			
	xor	%eax,%eax			
	nop				
5.64	200:	vmulsd	%xmm2,%xmm2,%xmm4		
5.47		vmovapd	%xmm2,%xmm3		
5.26		vaddsd	%xmm1,%xmm1,%xmm2		
5.54		lea	0x1(%rax),%rcx		
5.14		vfmsub231sd	%xmm1,%xmm1,%xmm4		
5.45		vfmadd213sd	%xmm6,%xmm3,%xmm2		
9.16		vaddsd	%xmm4,%xmm5,%xmm1		
9.30		vmulsd	%xmm2,%xmm2,%xmm3		
9.64		vfmadd231sd	%xmm1,%xmm1,%xmm3		
19.08		vucomisd	%xmm3,%xmm0		
9.05	↑	jb	1e0		
3.60		cmp	\$0x270f,%rax		
3.75		mov	%rcx,%rax		
3.91	↑	jb	200		
0.03		jmp	1e0		
23a:	mov	%rbx,%rdi			

No Abstraction

	data10	data10	CS	noplw	0x0(%rax)
1e0:	mov	%rcx,0x80(%rsp)			
	dec	%r14			
	je	238			
1ed:	vxorpd	%xmm1,%xmm1,%xmm1			
	vxorpd	%xmm2,%xmm2,%xmm2			
	vxorpd	%xmm3,%xmm3,%xmm3			
	vxorpd	%xmm4,%xmm4,%xmm4			
	xor	%eax,%eax			
	nop				
3.67	200:	vmovapd	%xmm4,%xmm5		
4.70		vaddsd	%xmm3,%xmm3,%xmm4		
6.82		vsubsd	%xmm1,%xmm2,%xmm1		
6.46		lea	0x1(%rax),%rcx		
7.78		vfmadd213sd	%xmm7,%xmm5,%xmm4		
10.05		vaddsd	%xmm1,%xmm6,%xmm3		
9.78		vmulsd	%xmm3,%xmm3,%xmm2		
9.87		vmulsd	%xmm4,%xmm4,%xmm1		
7.94		vaddsd	%xmm1,%xmm2,%xmm5		
16.23		vucomisd	%xmm5,%xmm0		
4.30	↑	jb	1e0		
4.66		cmp	\$0x270f,%rax		
3.39		mov	%rcx,%rax		
4.36	↑	jb	200		
		jmp	1e0		
238:	mov	%rbx,%rdi			

Computation Reuse

Assembly Comparison x86_64-v3

Key Observations

- Now the simpler code is faster, our computation reuse now hinders codegen
 - ▶ No abstraction: 1x move, 2x multiplications, 3x add/subs 3x fma
 - ▶ Computation reuse: 1x move, 2x multiplications, 5x add/subs, 1x fma
- Not obvious why the simpler code is faster
- Sometimes KISS (keep it simple & straightforward) is better

Pipeline Analysis with LLVM-MCA

What is LLVM-MCA?

- A performance analysis tool that simulates CPU pipeline behavior
- Models out-of-order execution, register renaming, and instruction scheduling
- Helps identify pipeline stalls, port pressure, and bottlenecks

Why Use It?

- Evidence driven approach to understand non-obvious performance differences
- Identifies resource contention and dependencies
- Helps understand the impact of instruction scheduling
- Particularly useful for analyzing microarchitecture-specific behavior

LLVM-MCA Analysis

Commands Used

```
clang++-20 mandelbrot.cpp -std=c++23 -O3 -S -o - |  
    llvm-mca-20 -mcpu=haswell
```

LLVM-MCA Analysis

Results Comparison

Metric	No Abstraction	Computation Reuse
Iterations	100	100
Instructions	1,900	2,100
Total Cycles	572	621
Total uOps	2,100	2,300
uOps/Cycle	3.67	3.70
IPC	3.32	3.38
Block RThroughput	5.3	5.8

Key Insight

The no-abstraction version has better (lower) Block RThroughput (5.3 vs 5.8), meaning it has better port scheduling and hence hides its latency better.

What about SIMD?

Single Instruction, Multiple Data

- A parallel computing technique where one instruction operates on multiple data elements simultaneously
- Modern CPUs have dedicated SIMD instruction sets (SSE, AVX, AVX2, AVX-512)
- Perfect for operations on arrays, vectors, and mathematical computations

Benefits for Mandelbrot Generation

- Calculate multiple results simultaneously
- Better utilization of CPU resources
- Mandelbrot is embarrassingly parallel

Scalar Addition (Traditional)

A:	1.0	2.0	3.0	4.0
B:	5.0	6.0	7.0	8.0
	+	+	+	+
C:	6.0	8.0	10.0	12.0

Scalar Processing

- 4 separate CPU instructions required
- Sequential execution
- One operation per clock cycle

SIMD Addition (AVX2)

A:

1.0	2.0	3.0	4.0
-----	-----	-----	-----

B:

5.0	6.0	7.0	8.0
-----	-----	-----	-----



Single SIMD instruction

C:

6.0	8.0	10.0	12.0
-----	-----	------	------

SIMD Processing

- Single CPU instruction (e.g., `vaddpd`)
- Parallel execution within one instruction
- 4x theoretical speedup for this example

SIMD Mandelbrot Implementation

xsimd-based Implementation

```
auto mandelbrot(xsimd::batch<double> a, xsimd::batch<double> b)
-> xsimd::batch<std::size_t> {
    using batch = xsimd::batch<double>;
    using bsize = xsimd::batch<std::size_t>;

    auto const four = batch(4.0);
    auto const two = batch(2.0);
    auto const one = bsize(1);

    auto x = batch(0.0);
    auto y = batch(0.0);
    auto iter = bsize(0);

    for (std::size_t i = 0; i < MAX_ITER; ++i) {
        auto const x2 = x * x;
        auto const y2 = y * y;

        auto const mask = (x2 + y2) <= four;
        if (none(mask)) {
            break;
        }

        auto const xy = x * y;
        auto const mask_i = batch_bool_cast<std::size_t>(mask);

        x = x2 - y2 + a;
        y = fma(two, xy, b);
        // Only update where still running
        iter = select(mask_i, iter + one, iter);
    }

    return iter;
}
```

SIMD Mandelbrot Key Concepts

Processing Multiple Complex Numbers

- `xsimd::batch<double>` gives largest size registers for your architecture, AVX2 would pack 4 doubles
- `mask = (x2 + y2) <= four` creates boolean mask for non-diverged points
- `select(mask, new_value, old_value)` conditionally updates only active points
- `if (none(mask))` early exit when all points have diverged

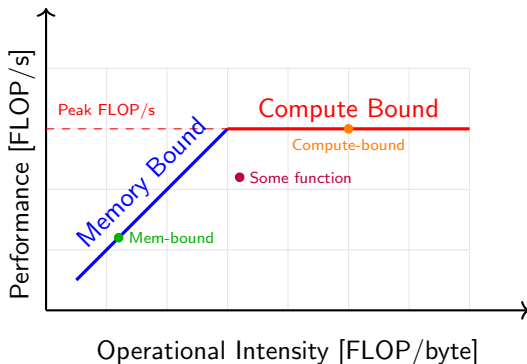
Performance Benefits

- 4x theoretical speedup from processing 4 complex numbers simultaneously
- Efficient handling of divergence without branching
- Modern CPU features: FMA, efficient mask operations

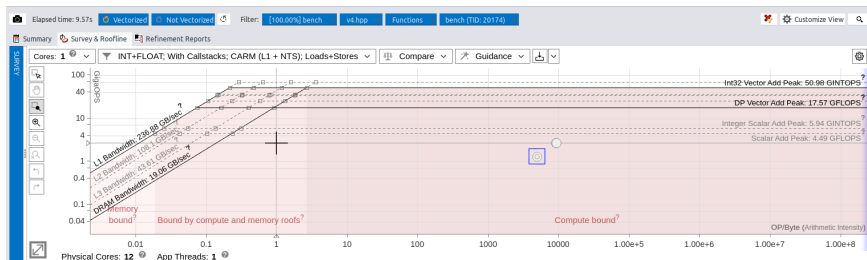
Intel Advisor for Performance Analysis

What is Intel Advisor?

- Performance profiling and optimization tool from Intel
- Provides vectorization analysis and roofline modeling
- Identifies optimization opportunities in CPU-bound applications



Compare scalar to SIMD



V4: No Abstraction (Scalar)



V6: SIMD Vectorization

Possible extra SIMD Optimizations choices

Loop Unrolling Strategies

- **Reduced escape checking:** Check every 4, 8, or 16 iterations
- **Compiler hints:** `#pragma clang loop unroll_count(16)`
- **Benefit:** Reduces branch prediction penalties

Is it worth it?

What is the likelihood of all members escaping? Are we ok to delay early exit by upto N-1 iterations?

SIMD with Loop Unrolling Implementation

unrolled xsimd-based Implementation

```
auto mandelbrot(xsimd::batch<double> a, xsimd::batch<double> b)
-> xsimd::batch<std::size_t> {
    using batch = xsimd::batch<double>;
    using bsize = xsimd::batch<std::size_t>;

    auto const four = batch(4.0);
    auto const two = batch(2.0);
    auto const one = bsize(1);

    auto x = batch(0.0);
    auto y = batch(0.0);
    auto iter = bsize(0);

#pragma clang loop unroll_count(16)
    for (std::size_t i = 0; i < MAX_ITER; ++i) {
        auto const x2 = x * x;
        auto const y2 = y * y;

        auto const mask = (x2 + y2) <= four;
        if (i % 16 == 0 and none(mask)) {
            break;
        }

        auto const xy = x * y;
        auto const mask_i = batch_bool_cast<std::size_t>(mask);

        x = x2 - y2 + a;
        y = fma(two, xy, b);
        // Only update where still running
        iter = select(mask_i, iter + one, iter);
    }

    return iter;
}
```

Modern CPUs: Beyond Single Core

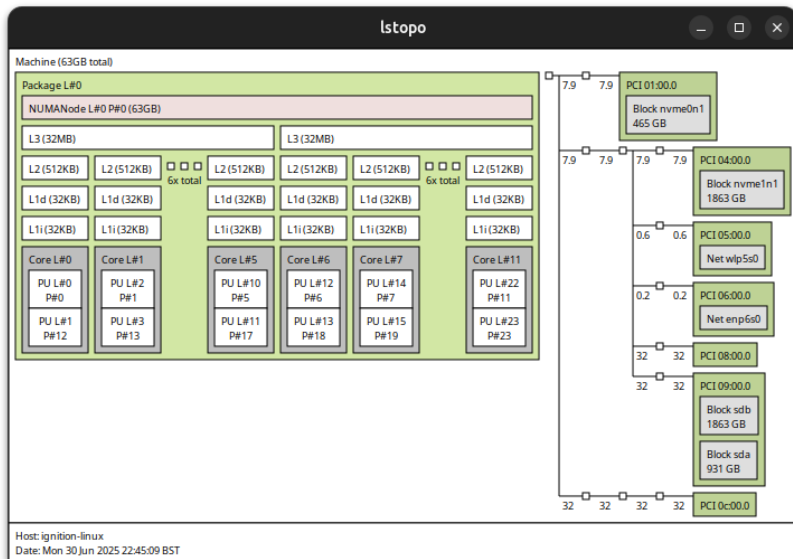
The Multi-Core Reality

- CPUs are no longer typically single core
- Even budget processors now have 4+ cores
- High-end consumer CPUs: 8-24+ cores
- Server CPUs: 64+ cores becoming common

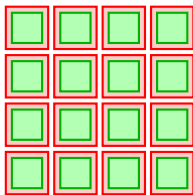
Why Single-Threaded Limits Performance

- Only utilizing 1/N of available CPU resources
- Mandelbrot calculation is embarrassingly parallel
- Each pixel calculation is independent
- Perfect candidate for multithreading

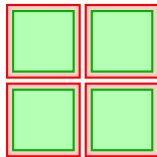
Hardware Topology



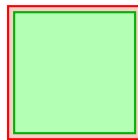
Concurrent Overhead Visualization



High Overhead



Low Overhead

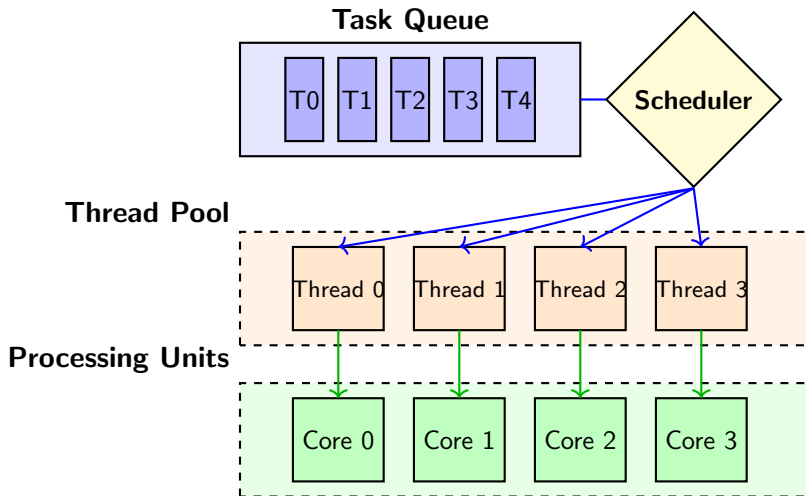


Low Concurrency

Overhead vs Useful Work

- **Red border:** Scheduling/synchronization/concurrency mechanism overhead
- **Green fill:** Useful computation work

Thread Pool Task Scheduling



Multithreaded SIMD Implementation

stdexec-based Implementation

```
void mandelbrot(std::vector<xsimd::batch<std::size_t>> &vec,  
               auto &&gen, auto scheduler) {  
    auto sender = stdexec::bulk(  
        stdexec::schedule(scheduler),  
        stdexec::par,  
        vec.size(),  
        [&](std::size_t i) {  
            vec[i] = std::apply(mandelbrot_simd, gen(i));  
        });  
    stdexec::sync_wait(sender);  
}
```

Key Components

- **stdexec::bulk**: Splits the work to be done as multiple chunked tasks
- **stdexec::par**: Parallel execution policy
- **scheduler**: Dispatch + coordinate via a thread pool
- **gen(i)**: Generates coordinates for batch i
- **mandelbrot_simd**: SIMD version from earlier
- **sync_wait**: Waits for all tasks to complete

Performance Results: Speedup vs Naive

Individual Optimizations

- Scalar optimizations: 1.8-2.2x improvement
- SIMD vectorization: 8-8.5x improvement
- Multithreading: 41x improvement (on 24 processing units)
- Combined MT + SIMD: **174x improvement**

The Journey: From Maths to Machines

Mathematical Foundation

- Started with elegant mathematical formula: $z_{n+1} = z_n^2 + c$
- Explored the complex plane and fractal geometry
- Translated mathematical concepts into computational algorithms

Tools Gave Us Insight

- **Profiling tools:** Hotspot, perf, Intel Advisor revealed bottlenecks
- **Assembly analysis:** Understanding what the compiler actually generates
- **LLVM-MCA:** Pipeline analysis for microarchitecture optimization
- **Benchmarking:** Evidence-driven development with Google Benchmark

Understanding Your Platform

Hardware Awareness Unlocks Performance

- **CPU microarchitecture:** x86_64-v3 enabled FMA and other better instructions
- **SIMD capabilities:** AVX2 provided 4x theoretical speedup, achieved 8.0x
- **Multi-core reality:** 24 processing units → 174x final speedup

Task Granularity is Critical

- Too fine: Overhead dominates useful work
- Too coarse: Poor load balancing and resource waste
- Sweet spot: Balance between parallelism and scheduling costs

The Learning Never Stops

Technology Evolves Continuously

- **Hardware changes:** New instruction sets, architectures, and capabilities
- **Languages evolve:** C++26 brings new abstractions (sender/receivers) [libstdexec]
- **Libraries advance:** xsimd provides portable SIMD, new algorithms emerge

Fundamental Principles Remain

- Measure first, optimize second
- Understand your problem domain and hardware platform
- Balance abstraction with performance requirements
- Use tools to guide decisions, not assumptions

Final Thoughts

174x Speedup: The Journey Matters

- Mathematical elegance → Clean initial implementation
- Profiling insights → Targeted optimizations (sqrt → norm)
- Assembly understanding → Manual improvements
- Hardware awareness → SIMD vectorization
- Platform knowledge → Effective multithreading

Keep Exploring

- Try other fractal sets (Julia, Burning Ship, Newton)
- Experiment with GPU computing (CUDA, OpenCL, compute shaders)
- Explore distributed computing and cloud scaling
- Different coloring techniques
- Getting past IEEE double precision limits